



Apples-to-apples comparison on the digital-twin 70/30 split (n $\approx$ 302-303 per domain)

Aggregate accuracy: comparable



Within-participant rank discrimination: LLM only



Supervised RF emits $\sim$ constant predictions across messages within a participant; rank ρ is therefore degenerate (≈ 0).