

Apples-to-apples scaling: supervised RF versus LLM-PP (k_{train} in $\{1, 3, 7\}$)

